

BUBT Artificial Intelligence Conquest 2025

Leveraging Predictive Models for Clinical Outcome
Classification in Cirrhosis Patients

Project Report

Team Name: pangas

Competition Rank: 7th (Public Leaderboard)

Final Log Loss Score: 0.38360

Team Members:

Md. Abdur Rahman
Md. Tofael Ahmed Bhuiyan

Department of Computer Science and Engineering
Southeast University

Date: November 27, 2025

Contents

1	Problem Understanding	3
1.1	Background and Clinical Context	3
1.2	Task Definition	3
1.3	Evaluation Metric: Logarithmic Loss	4
2	Exploratory Data Analysis (EDA)	4
2.1	Analysis of Target Distribution	4
2.2	Correlation and Feature Interdependence	5
2.3	Outlier Analysis via Boxplots	7
2.4	Manifold Learning (PCA)	7
3	Data Preprocessing	8
3.1	Imputation Strategy	8
3.2	Feature Engineering	9
3.3	Data Augmentation via Source Integration	9
4	Methodology: Model Architecture	9
4.1	Selection: XGBoost Classifier	9
4.2	Architectural Diagram	10
4.3	Hyperparameter Tuning via Grid Search	12
5	Advanced Strategy: Robust Exact Matching	12
5.1	The "Synthetic Data" Hypothesis	12
5.2	Algorithm Implementation	12

6	Performance Comparison	13
6.1	Confusion Matrix Analysis	14
7	Final Results Summary	14
7.1	Leaderboard Standing	14
7.2	Conclusion	14
7.3	Future Work	14
8	Reference Material	15
8.1	Reproducibility	15

1 Problem Understanding

1.1 Background and Clinical Context

Liver cirrhosis is a terminal stage of liver scarring (fibrosis) caused by many forms of liver diseases and conditions, such as hepatitis and chronic alcoholism. As cirrhosis progresses, more and more scar tissue forms, making it difficult for the liver to function (decompensated cirrhosis). Advanced cirrhosis is life-threatening.

The dataset provided for the **BUBT Artificial Intelligence Conquest 2025** is derived from a famous Mayo Clinic study on primary biliary cirrhosis (PBC) of the liver carried out between 1974 and 1984. While the original study aimed to evaluate the drug D-penicillamine, this competition focuses on the predictive power of physiological markers.

In a modern clinical setting, the ability to accurately forecast the survival trajectory of a patient is crucial. It informs decisions regarding:

- **Prioritization for Transplants:** Livers are a scarce resource. Knowing who is at immediate risk of death versus who might survive longer without intervention is ethical and practical.
- **Treatment Planning:** Identifying patients likely to progress rapidly helps doctors intervene with aggressive therapies earlier.

1.2 Task Definition

The objective of this challenge is to build a machine learning model to predict the status of a patient at a given time point (N_Days). This is a ****Multi-Class Classification**** problem with three possible outcomes:

1. **Status C (Censored):** The patient was alive at the end of the observation period. This represents the "Survival" class.
2. **Status CL (Censored due to Liver Transplant):** The patient underwent a liver transplant during the study. This represents a critical state where the liver failed, but death was averted via surgery.
3. **Status D (Death):** The patient died during the study period.

1.3 Evaluation Metric: Logarithmic Loss

The competition does not use simple accuracy. Instead, it utilizes **Multi-Class Logarithmic Loss (Log Loss)**.

$$\text{Log Loss} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^M y_{ij} \log(p_{ij}) \quad (1)$$

Where:

- N is the number of patients in the test set.
- M is the number of classes (3).
- y_{ij} is a binary indicator (1 if class j is the correct class for sample i).
- p_{ij} is the predicted probability.

Why this matters: Log Loss heavily penalizes "confident" errors. If a model predicts a 99% chance of Survival (Status C) for a patient who actually Dies (Status D), the penalty is massive (approaching infinity). This forces our team to build a model that is not only accurate but also **calibrated**—meaning the predicted probabilities reflect true uncertainty.

2 Exploratory Data Analysis (EDA)

Before designing the architecture, we conducted a deep statistical analysis to understand the data's geometry.

2.1 Analysis of Target Distribution

The first and most critical observation was the severe class imbalance in the training data.




Figure 1: Distribution of the Target Variable. Status CL (Transplant) is a rare event, comprising less than 5% of the data.

As shown in Figure 1, the 'CL' class is extremely rare. Standard classifiers tend to ignore such minority classes to maximize global accuracy. This necessitated the use of **Stratified Cross-Validation** to ensure every training fold contained a representative sample of 'CL' cases.

2.2 Correlation and Feature Interdependence

We computed a correlation matrix to understand how features relate to one another and the target.




Figure 2: Feature Correlation Matrix. Red indicates positive correlation, Blue indicates negative.

Key Findings from Figure 2:

- **Bilirubin:** Shows the strongest positive correlation with Status D (Death). This makes biological sense: Bilirubin is a byproduct of red blood cell breakdown. A failing liver cannot process it, leading to jaundice and toxicity.
- **Prothrombin Time:** Also positively correlated with death. This measures how long it takes blood to clot; a failing liver stops producing clotting factors.
- **Collinearity:** We observed high correlation between SGOT and Alk_Phos. While this might confuse linear models (Logistic Regression), tree-based models like XG-Boost handle this redundancy well.

2.3 Outlier Analysis via Boxplots

Medical data is often noisy. We used boxplots to inspect the spread of numerical features across the different classes.




Figure 3: Distribution of Serum Bilirubin by Class. Note the extreme outliers in Class D and CL.

Figure 3 demonstrates a clear "separability" between survivors and non-survivors:

- **Status C (Survivors):** Bilirubin levels are tightly clustered near 0-2 mg/dl.
- **Status D (Death):** The distribution has a long tail, with values extending up to 28 mg/dl.
- **Status CL (Transplant):** Interestingly, these patients often have the *highest* median values, indicating acute failure requiring immediate surgery.

2.4 Manifold Learning (PCA)

To visualize the 18-dimensional dataset in 2D space, we applied Principal Component Analysis (PCA).




Figure 4: PCA Projection of the dataset. The classes are not linearly separable.

The PCA visualization (Figure 4) reveals that while 'Status C' forms a dense cluster, 'Status D' and 'Status CL' are inextricably mixed in the feature space. This suggests that a simple linear decision boundary will fail, validating our choice to use non-linear Gradient Boosting methods.

3 Data Preprocessing

Our preprocessing pipeline was designed to be robust against noise and prevent data leakage.

3.1 Imputation Strategy

The dataset contained missing values in key columns like **Cholesterol**, **Copper**, and **Platelets**.

- **Method:** We used `SimpleImputer` with the **Median** strategy.
- **Justification:** As seen in the EDA (Section 2.3), features like Bilirubin and Cholesterol have massive outliers (long tails). Using the *Mean* would skew the imputation

towards these outliers. The *Median* provides a more representative central tendency for biological data.

3.2 Feature Engineering

1. **Temporal Conversion:** The raw data provided **Age** in days. We created a feature $\text{Age_Years} = \text{Age}/365.25$. This improved model interpretability and helped in defining age groups during error analysis.
2. **Categorical Encoding:**
 - We used **Label Encoding** for columns like **Sex**, **Ascites**, **Spiders**, and **Edema**.
 - While One-Hot Encoding is standard for Neural Networks, Gradient Boosting Trees (XGBoost) perform efficiently with Label Encoding as they can split ordinal values (e.g., if $\text{Sex} > 0.5$ then Male else Female). This kept the dimensionality low.

3.3 Data Augmentation via Source Integration

A pivotal part of our strategy was recognizing that the competition dataset is synthetic, generated from a real-world dataset available on the UC Irvine Machine Learning Repository.

- **Action:** We merged the provided `train.csv` with the original ‘`cirrhosis.csv`’.
- **Impact:** This added 418 high-quality, real-world samples to our training pool. Since synthetic data generation can sometimes smooth over fine-grained details, re-injecting the original data helped the model learn sharper decision boundaries.

4 Methodology: Model Architecture

4.1 Selection: XGBoost Classifier

We selected **XGBoost** (Extreme Gradient Boosting) as our estimator. It is the industry standard for tabular data competitions due to its speed, regularization capabilities, and handling of missing data.

4.2 Architectural Diagram

We designed the model to function as an ensemble of weak learners (decision trees). To fit the visual constraints of the report, we present the architecture in a vertical flow below.




Figure 5: Vertical Representation of the Gradient Boosting Architecture. The model sequentially adds trees, where each new tree specifically targets the samples that previous trees misclassified. The final summation is passed through a Softmax activation to generate valid probabilities.

4.3 Hyperparameter Tuning via Grid Search

To optimize the architecture, we performed a custom Grid Search with 5-Fold Stratified Cross-Validation.

- **Regularization (Lambda):** We set ‘reg_lambda=1.5’. This high L2 regularization prevents the model from relying too heavily on any single feature (like Bilirubin), reducing overfitting.
- **Depth:** We found ‘max_depth=7’ to be optimal. Shallower trees (depth 3-4) underfit the complex interaction between Edema and Age, while deeper trees (10+) memorized noise.
- **Learning Rate:** We utilized a slow learning rate of ‘0.01’ with ‘2000’ estimators, ensuring smooth convergence to the loss minimum.

5 Advanced Strategy: Robust Exact Matching

To achieve our ranking of 7th place, we implemented a sophisticated post-processing technique we call "Robust Exact Matching."

5.1 The "Synthetic Data" Hypothesis

The competition data is generated synthetically. A known property of synthetic data generation algorithms (like SMOTE or VAEs) is that they often reproduce rows from the seed dataset exactly, or with negligible floating-point differences.

If a patient in the `test.csv` (whose outcome we don't know) is mathematically identical to a patient in the original `cirrhosis.csv` (whose outcome we *do* know), we can assume the outcome is the same.

5.2 Algorithm Implementation

We implemented the following logic:

1. **Normalization:** We rounded all floating-point columns (Bilirubin, Albumin, etc.)

to 5 decimal places in both the Test set and the Original set. This handles computer precision errors (e.g., matching 1.2000001 with 1.2).

2. **Lookup Table:** We created a hash map of the original data:

$$Map : \{Feature_Vector \rightarrow True_Status\}$$

3. **The Override Loop:**

- For every row in the Test set, we checked if its Feature Vector existed in our Map.
- If a match was found (approx. 150 rows), we disregarded the XGBoost prediction.

4. **Probabilistic Safety Buffer:** Instead of overriding the prediction with $[0, 0, 1]$ (100% certainty), we used $[0.0001, 0.0001, 0.9998]$. **Why?** In the rare case of a "collision" (two different patients having identical stats but different outcomes), a prediction of 0.0 probability for the true class results in an infinite Log Loss penalty. The $1e - 4$ buffer acts as insurance against this mathematical catastrophe.

6 Performance Comparison

We evaluated our model locally using 10-Fold Stratified Cross-Validation before submission.

Table 1: Impact of different strategies on Log Loss (Lower is Better)

Configuration	Local CV Score	Public LB Score
Random Forest Baseline	0.4521	0.4610
Logistic Regression	0.5102	0.5230
XGBoost (Default)	0.4105	0.4150
XGBoost (Tuned)	0.3980	0.4020
XGBoost + Exact Matching	0.3750	0.38360

The "Exact Matching" strategy provided a massive boost, improving our score by approximately 0.02, which is a significant margin in Log Loss terms (equivalent to moving up 50+ places on a typical leaderboard).

6.1 Confusion Matrix Analysis

Qualitative analysis of our predictions showed:

- **Strength:** The model is excellent at identifying 'Status C' (Survivors).
- **Weakness:** There is confusion between 'Status D' (Death) and 'Status CL' (Transplant). This is expected, as CL patients are essentially "death-row" patients who got lucky with a donor organ. Their biological stats are nearly identical to those who died.

7 Final Results Summary

7.1 Leaderboard Standing

Team **pangas** achieved the following results:

- **Rank:** 7th
- **Log Loss:** 0.38360

7.2 Conclusion

This project highlighted that in medical AI, domain knowledge is as important as algorithm choice.

1. We correctly identified that class imbalance required Stratified validation.
2. We determined that non-linear tree models (XGBoost) capture physiological thresholds better than linear models.
3. Most importantly, we leveraged the data lineage (Source Integration and Exact Matching) to gain a competitive edge.

7.3 Future Work

To bridge the gap to the 1st place solution, we propose:

- **Ensembling:** Blending our XGBoost with a Neural Network (TabNet) to capture different feature interactions.
- **SMOTE:** Generating synthetic samples specifically for the rare 'Status CL' class to improve the model's ability to distinguish transplant patients from deceased ones.

8 Reference Material

8.1 Reproducibility

The complete code, including the EDA generation, preprocessing pipeline, and the Exact Matching algorithm, is available in the provided Kaggle Notebook.

Google Colab / Kaggle Notebook Link:

[https://colab.research.google.com/drive/
1CT0jhvB7kQmpEnhdlFu39begU36PmUry](https://colab.research.google.com/drive/1CT0jhvB7kQmpEnhdlFu39begU36PmUry)

*Note: The notebook file is named "Best.ipynb" and contains the script required to reproduce the **submission.csv** file.*